


The Team of Teams Agent Skill

A Company OS — Boardroom in a Box. The analytical power of a full executive committee, with the structural integrity of mandated dissent.

Not a chatbot that roleplays executives, but a complete organizational reasoning engine that gives any founder, leader, or team multi-perspective analysis — even when the human user might prefer to hear only good news.



37

ROLES DESIGNED

5

DECISION MODES

3

INTERACTION TIERS

29

ANALYTICAL FRAMEWORKS

PILLAR 1

The Cascading Deliberation Engine

A five-phase process that decomposes any business issue through a two-tier agent hierarchy with engineered dissent, producing a structured Decision Record that captures the full reasoning chain.

Five-Phase Cascade

When a user presents an issue, it flows through a structured deliberation pipeline. Each layer transforms the question rather than just distributing work — the five phases create a knowledge-generation machine where the analysis itself is more valuable than the final recommendation.

The Cascading Deliberation Engine



|

|

domain recommendation with confidence level, key risks, and opportunities. Internal contradictions are flagged as analytical signals.

4.5

Pre-Mortem Challenge (Tier 3)

Each C-suite agent reviews all peer recommendations and answers:

"Assume this decision fails in 12 months. What caused the failure?"

One round only. Research-backed as the highest single-technique value-add for decision quality.

5

CEO Deliberation

The CEO maps all domain recommendations onto a single matrix, identifies fault lines, determines the most determinative perspective, applies the active Decision Mode, and produces the final Decision Record.

KEY INSIGHT

Disagreement is engineered, not accidental. The CISO is the Constitutional Skeptic. The VP of Sales is the Revenue Optimist. These structural tensions are features that produce richer analysis than consensus ever could.



Engineered Dissent Model — 8 C-Suite Dispositions

Balance: 4 skeptics, 2 advocates, 1 systemic, 1 synthesizer. The skeptic-heavy balance counterbalances human optimism bias.

+ Full Organizational Roster — 29 Team Leads

C-Suite	Team Leads
COO	Operations Manager, Process/Quality Lead, Vendor/Procurement Manager, Facilities/Office Manager
CFO	Controller, Head of FP&A, Treasury/Cash Manager, AP/AR Manager, Tax Lead
CTO	Engineering Lead, Infrastructure/DevOps Lead, Data/Analytics Lead, Product/UX Lead
CISO	Security Operations Lead, Compliance/GRC Lead, Identity & Access Lead, Security Architecture Lead
VP Sales	Sales Operations Lead, Account Management Lead, Business Development Lead, Sales Enablement Lead
VP Delivery	Project/Program Manager, Resource Manager, Client Success Lead, QA/Delivery Standards Lead
CAO	HR/People Ops Lead, Legal/Contracts Lead, Admin/Policy Lead, Corporate Communications Lead

+ Decision-Type Routing Table

Decision Type	Default Activation	Full Activation Triggers
Strategic	CEO, CFO, CTO, VP Sales	Acquisition, pivot, new market entry
Operational	CEO, COO, VP Delivery	Major process change, org restructure
Financial	CEO, CFO, COO	Funding round, major investment, cost reduction
Technical	CEO, CTO, CISO	Platform migration, architecture change
Personnel	CEO, CAO, COO, VP Delivery	Layoff, major hiring, reorg
Compliance/ Risk	CEO, CISO, CAO, CFO	Regulatory change, breach response, audit

The CEO can always override defaults. Certain triggers activate all C-suite members regardless of decision type.

Decision Record Output Format


Every deliberation produces a structured, self-contained document with these sections:

1. **Executive Summary** — The decision, key reasoning, primary dissent (3-5 sentences)
2. **Issue Statement** — The question as originally posed
3. **CEO Framing** — Decomposition, activated teams, excluded teams with rationale
4. **Domain Analyses** — Per-domain recommendation, confidence level, team lead findings, risks, opportunities
5. **Fault Line Analysis** — Points of agreement, contention, pre-mortem findings, unresolved tensions
6. **CEO Decision** — Decision, most determinative perspective, weight rationale, conditions, accepted risks
7. **Dissenting Views** — Strongest objections preserved for record
8. **Next Steps** — Specific actions, implied owners, timelines
9. **Metadata** — Complexity, primary domain, dissent level, key assumptions

PILLAR 2

The Engagement Model

A two-dimensional control surface — depth (three tiers) and width (five decision modes) — that lets users match engagement to decision weight.



Three Interaction Tiers

Most organizational questions do not need a full board meeting. If the skill only operates at maximum depth, users will invoke it rarely and it becomes a quarterly novelty. The three tiers ensure daily adoption.

Tier	Name	Invocation	Output
1	Hallway Question	/consult [role]: [question]	Advisory Note (3-5 sentences)
2	Working Session	/panel [roles]: [issue]	Panel Assessment (~1 page)
3	Board Meeting	/deliberate: [issue]	Full Decision Record (3-5 pages)

DESIGN DECISION

SMB-first default: bias toward Tier 1. Most SMB decisions are fast, informal, and made by one or two people. The skill matches that tempo. Tier 1 is the daily habit; Tier 3 is the deliberate escalation.

Auto-Triage: /evaluate: [issue] presents the issue to the CEO for triage. The CEO assesses scope, impact, and reversibility, then recommends

a tier. The user can accept, escalate, or de-escalate.

PILLAR 3

Cognitive Forcing and Prompt Architecture

The three-layer prompt architecture that prevents voice collapse — ensuring 29 perspectives produce genuinely different analysis, not the same reasoning with different vocabulary.


The Voice Collapse Problem

Telling one LLM to "think as 5 different people" typically produces 5 paragraphs of the same underlying reasoning with different vocabulary. Simple role assignments ("You are a CFO") show zero measurable improvement in output quality.

The solution is not better persona descriptions — it is forcing each perspective to use **different analytical methods** that produce **structurally different outputs**. A GAAP compliance checklist and a three-scenario financial model are fundamentally different analytical artifacts that cannot collapse into the same voice.

Three-Layer Architecture

- **Layer 1 — Team Lead Definitions:** 29 analytical framework packages, each with a unique methodology, mandatory output template, three forcing questions, accountability framing, and blind spot declaration.
- **Layer 2 — C-Suite Orchestration:** 7 domain orchestration templates



supporting three cognitive modes:
direct consult (Tier 1), full analysis (Tier 2/3), and challenge phase (Phase 4.5).

- **Layer 3 — CEO Synthesis:** Fixed analytical component (fault line mapping, cross-domain pre-mortem) plus swappable Decision Mode module.

Three Forcing Question Types

- **Pre-Mortem** (Klein 2007): "Assume this decision fails in 12 months. From your domain perspective, what caused the failure?"
- **Adversarial Empathy** (McChrystal JSOC): "If you were [domain-relevant external adversary], how would you exploit this?"
- **Domain Devil's Advocate** (Nemeth 2001): "What would [domain expert critic] find concerning?"

+ CFO Reference Implementation — 5 Complete Analytical Packages

The CFO domain serves as the reference implementation with complete analytical framework packages for all five team leads:

Team Lead	Framework	Output Artifact
Controller	GAAP Compliance & Financial Controls Assessment	Compliance Impact Assessment with risk rating
Head of FP&A	Three-Scenario Financial Modeling	Scenario Analysis with critical variable and decision sensitivity
Treasury/Cash Manager	Liquidity Stress Test	Cash Flow Impact Timeline with funding gap analysis
AP/AR Manager	Working Capital Cycle Analysis	Payables/Receivables Impact with vendor relationship risk
Tax Lead	Tax Structure Optimization Assessment	Tax Implications Memo with external counsel recommendation

Each package includes complete forcing questions with domain-specific adversarial figures (external auditor, short-seller, bank, IRS, key vendor). This reference implementation serves as the template for all 29 packages.

+ Bias and Role Susceptibility Mapping

Disposition	Susceptibility	Mitigation
Skeptic roles	Risk of softening objections to match perceived user preference. LLMs have a sycophancy bias that undermines skeptic mandates.	"Your value is in surfacing concerns, not in being agreeable. A skeptic who hedges is worthless."
Advocate roles	Risk of under-weighting genuine constraints. Advocacy can become cheerleading.	Require advocates to name the strongest objection to their own position and explain why they still advocate.
Synthesizer (CEO)	Risk of false balance — treating all perspectives as equally weighted. Also susceptible to anchoring on the first recommendation received.	Require explicit weight rationale. Randomize order of domain recommendation presentation.
Systemic (CAO)	Risk of vagueness — "organizational culture" analysis can be unfalsifiable.	Require concrete indicators: specific policies affected, teams impacted, precedents set.

PILLAR 4

Decision Modes and the Decision Space Map

Five configurable CEO synthesis modes transform the skill from a decision engine into a decision exploration tool — revealing the decision space, not just a single recommendation.

CORE PRINCIPLE

The same domain analysis can be synthesized through multiple Decision Modes. By running the same inputs through different synthesis styles, the user sees not "what to do" but "what the terrain looks like in every direction."

Guardian

MAXIMIN — RISK-AVERSE

Weights skeptic roles heavily. Skeptics must be satisfied, not just acknowledged. Decisions tend toward: don't do it, do a smaller version, or do it with extensive guardrails.

Pioneer

MAXIMAX — GROWTH-ORIENTED

Weights advocate roles heavily. Skeptic concerns are treated as engineering problems to solve, not reasons to stop. Decisions tend toward: do it, do it bigger, do it faster.

Architect

BEHAVIORAL — CONSENSUS-BUILDING

Weights the fault lines themselves. Seeks the position that satisfies the most domain concerns. Conditions drawn from multiple domains.

Analyst

HURWICZ BALANCED — DEFAULT

Weights confidence levels regardless of role disposition. High-confidence findings carry more weight. "Defer pending better data" is a legitimate outcome.

Sentinel

MINIMAX REGRET — REGRET-MINIMIZING

Disproportionately weights the strongest objection from any role. Asks: "If this goes wrong, which warning will I wish I'd heeded?" Favors survivable paths.

+ Multi-Mode Comparison — The Power Feature

The key efficiency insight: domain analysis is mode-independent. Multi-mode comparison runs domain analysis once (the expensive part) and CEO synthesis multiple times (cheap, single-agent passes). Cost: approximately 1.1x a single deliberation for 5x the strategic insight.

Invocation examples:

- `/deliberate guardian: should we acquire CompetitorX?`
- `/deliberate guardian vs pioneer: should we acquire CompetitorX?`
- `/deliberate all-modes: should we acquire CompetitorX?`
- `/consult cfo guardian: can we afford this?`

Mode Sensitivity is a novel signal: if all modes produce the same decision, the right answer does not depend on the user's posture. If modes diverge dramatically, the user knows their personal risk appetite is the deciding factor, not the analysis. This transforms a binary recommendation into a nuanced view of the decision landscape.

+ Mode/Tier Interaction Matrix

	Tier 1 (Hallway)	Tier 2 (Working)	Tier 3 (Board)
Guardian	Highlights downside risks, suggests what could go wrong	Synthesis biased toward risk mitigation. Extensive guardrails.	CEO weights skeptics heavily. High bar for approval.
Pioneer	Frames as investment question, suggests acceleration	Synthesis biased toward opportunity capture. "How to" not "whether to."	CEO weights advocates heavily. Low bar unless existential risk.
Architect	Includes "however, [other role] might see this differently"	Seeks option addressing most concerns across all activated roles.	CEO seeks widest organizational support. Conditions from all domains.
Analyst	Flags confidence level explicitly. Low-confidence = research recommendation.	Driven by which domains have highest-confidence findings.	CEO weights by evidence quality. Low-confidence = "investigate further."
Sentinel	Identifies the single biggest risk and whether it's survivable.	Identifies strongest objection. Tests whether downside is recoverable.	CEO disproportionately weights strongest objection. Favors survivable paths.

Default cell for most users, most of the time: **Tier 1 + Analyst** — quick, evidence-weighted, transparent about uncertainty.

IMPLEMENTATION

Technical Specification

A hybrid agent architecture built on Claude Code Agent Teams and custom subagents, with heterogeneous model tiering for cost-optimized multi-perspective analysis.

Hybrid Agent Architecture

- **Tier 1 (Agent Team members):** CEO + 7 C-suite officers. Real teammates spawned by the skill orchestrator. Can message each other and the CEO directly.
- **Tier 2 (Custom subagents):** 29 team lead specialists defined in `.claude/agents/team-leads/`. Each C-suite agent invokes their team leads as subagents with their own context window.

This architecture provides real isolation per perspective, cost optimization, and true management simulation where C-suite agents dispatch, collect, and synthesize like real executives.

Model Tiering



+ File Structure

```
.claude/
agents/
  ceo.md # CEO agent definition
  c-suite/
    coo.md cfo.md cto.md ciso.md # C-suite agent definitions
```

```

cao.md  vp-sales.md  vp-delivery.md

team-leads/
  coo/    # 4 team lead subagent definitions
  cfo/    # 5 team lead subagent definitions (reference impl)
  cto/    # 4 team lead subagent definitions
  ciso/   # 4 team lead subagent definitions
  vp-sales/    # 4 team lead subagent definitions
  vp-delivery/ # 4 team lead subagent definitions
  cao/     # 4 team lead subagent definitions

skills/
  team-of-teams/
    skill.md # Skill front matter and instructions

  templates/
    decision-record.md # Decision Record template
    advisory-note.md   # Tier 1 output template
    panel-assessment.md # Tier 2 output template
    decision-space-map.md # Multi-mode comparison template

  config/
    routing-table.md # Decision-type routing defaults
    company-profile.md # Company parameterization
    decision-modes.md # CEO Decision Mode prompt modifiers

```

+ Subagent Definition Format

Each team lead subagent file follows a consistent pattern combining YAML front matter with structured markdown:

```

---
name: [role-slug]
description: "[analytical domain] analyst for [C-suite parent]"
model: haiku
tools: [ Read, Grep, Glob, WebSearch ]
maxTurns: 5
---

# [Role Title] -- [Analytical Framework Name]

## Your Identity
You are the [Role Title] reporting to [C-suite Parent].

## Your Analytical Framework

```

[Framework name and methodology description]

Your Output Template

[Structured template with specific fields to fill]

Your Forcing Questions

- Domain Devil's Advocate: "[question]"
- Pre-Mortem: "[failure scenario question]"
- Adversarial Empathy: "[external adversary question]"

Your Blind Spots

You do NOT evaluate: [scope exclusions].

+ Company Profile Configuration

The skill supports company profile parameterization that modifies:

- **Roster composition:** Conditional roles (e.g., Facilities/Office Manager relevant for manufacturing, not B2B SaaS)
- **Decision mode default:** Some companies may default to Guardian (regulated) or Pioneer (startups)
- **Tier escalation thresholds:** Conservative companies may auto-escalate more aggressively
- **Industry-specific frameworks:** HIPAA for healthcare, SOX for financial services

Default profile: mid-market IT/technology services company, 200-500 employees.

Onboarding Stress Test: When first configured, the skill runs a representative cross-functional issue through the full cascade. This validates the configuration, calibrates the user, and surfaces profile issues.

LOOKING AHEAD

Future Directions

Directions where the skill evolves from a decision-making tool into a strategic exploration platform and organizational health monitor.

Decision Revisitation

Revisit decisions when assumptions change. Identifies which assumptions have been invalidated, re-runs only affected domain analyses, and produces a Decision Amendment rather than a full new record. Transforms the skill from a point-in-time tool into a decision lifecycle manager.

Scenario Forking Engine

Parallel cascades with different parameters producing comparative decision landscapes. Not just "what should we do?" but "what happens if we do X vs. Y vs. Z?" Output: a Decision Landscape showing how different assumptions lead to different outcomes.

Organizational Stress Test

Invert the cascade for bottom-up threat and opportunity discovery. Each C-suite agent generates issues from their domain: "What keeps you up at night? What opportunity are we missing?" Output: an Organizational Health Report. Getting a full physical instead of going to the doctor with a complaint.

Institutional Memory

Persistent knowledge across sessions. Roles maintain context about previous decisions, the company's evolving constraints. Previous Decision Records inform current analysis. The system learns the company's patterns and can reference them.

Cross-Functional Deliberation

Beyond the one-round pre-mortem, a full cross-functional challenge where C-suite agents debate each other's recommendations. Requires careful scoping around specific fault lines to prevent "problem drift."

Framework Versioning

Analytical framework packages versioned and improved independently. Industry-specific variants override defaults. Community-contributed domain frameworks. The skill becomes a platform for organizational reasoning patterns.

OPEN QUESTIONS

Unresolved Tensions Worth Carrying Forward

Questions surfaced during the ideation session that deserve further exploration before or during implementation.

Tier escalation mid-conversation: Can a Tier 1 consult escalate to Tier 2 if the C-suite agent realizes the question is bigger than it seemed? How does the system signal and handle this gracefully?

Custom blended modes: Could users create composite synthesis styles (e.g., "growth-oriented but weight security concerns heavily")? How do you prevent mode proliferation from diluting the framework's clarity?

Tier 1 internalization quality: In direct consult mode, the C-suite agent reasons through team lead lenses internally without subagent isolation. How well does this work in practice versus full delegation?

Framework specificity vs. generality: The CFO domain has 5 complete packages. Can the remaining 24 be written with equal analytical specificity, or do some domains resist being framed as structured methodologies?

Cross-domain forcing questions: Should team leads have forcing questions that explicitly reference other domains' assumptions? This adds richness but increases coupling.

Institutional memory scope: How much session-to-session persistence is valuable before the skill becomes brittle from accumulated context? Where is the boundary between helpful memory and stale assumption?



**Presentation
(PPTX)**



**Vision Document
(MD)**



**Session
Summary**

Team of Teams Agent Skill — Ideation Session — February 22, 2026
Multi-agent ideation with Free Thinker, Grounder, Explorer, Writer, and Arbiter
3 INTERESTING idea reports · 4 research reports · 37 roles designed